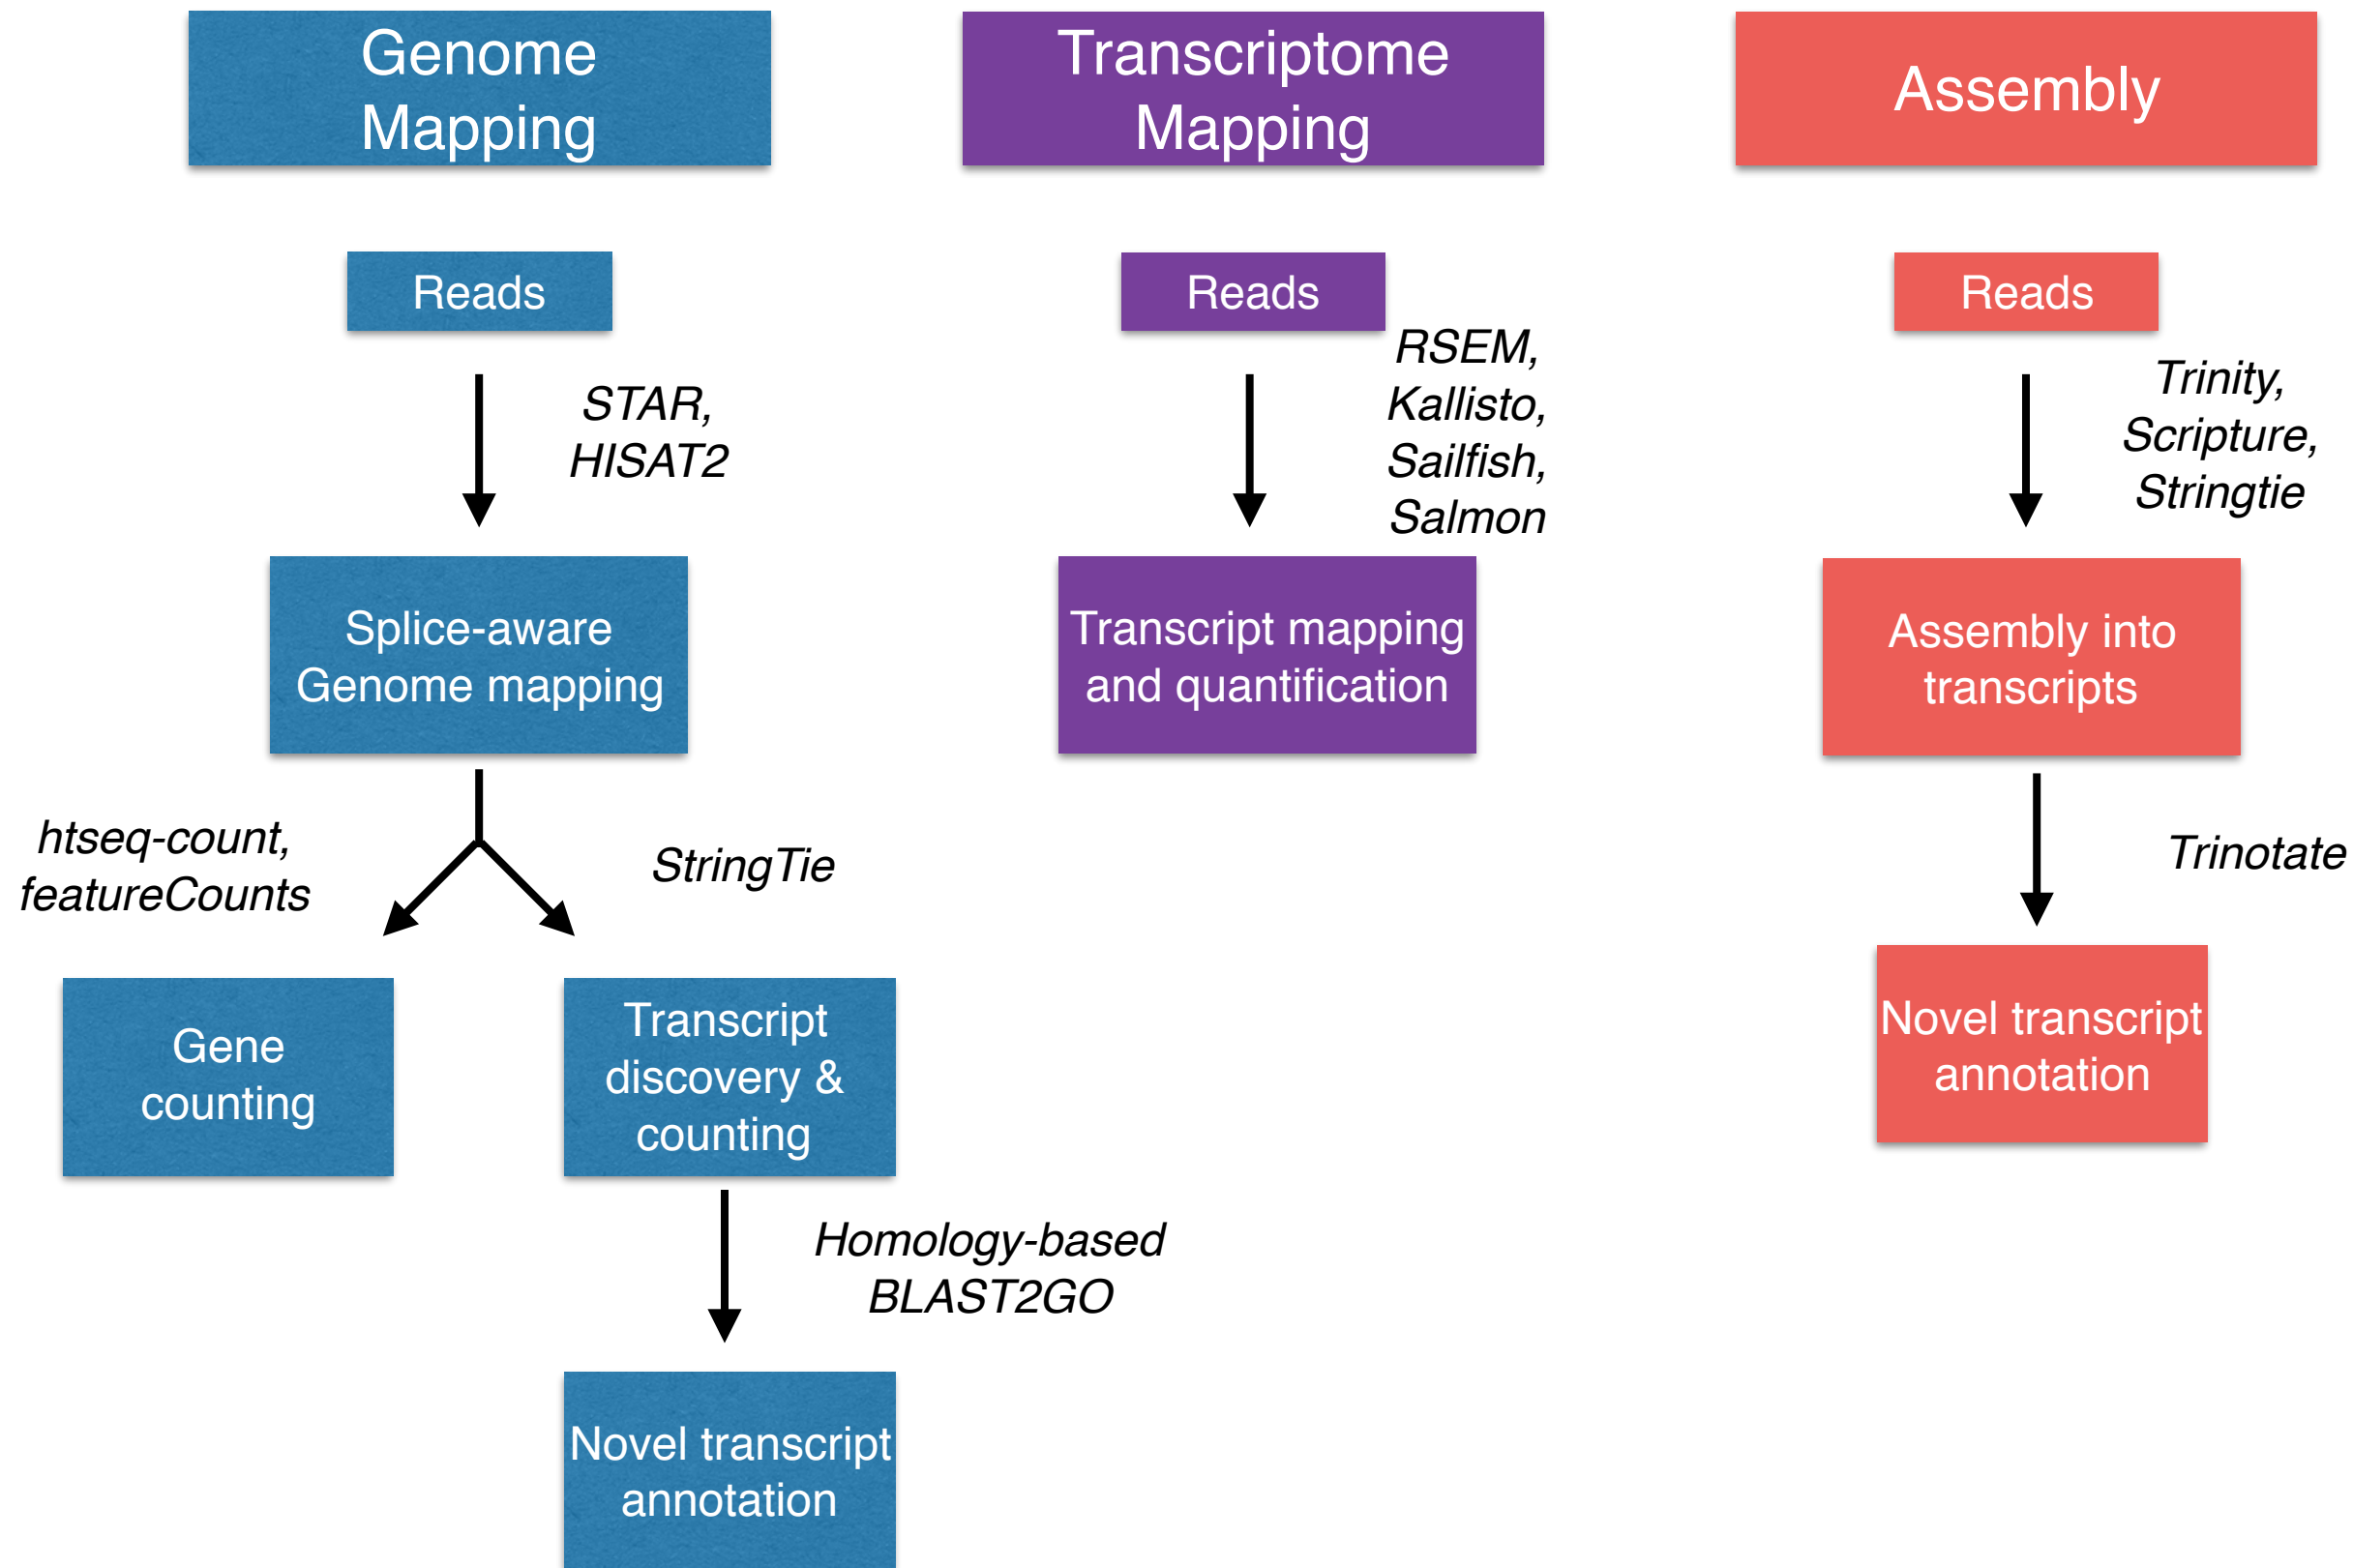


Aligning reads: tools and theory



Transcriptome Mapping

Reads



*RSEM,
Kallisto,
Sailfish,
Salmon*

Transcript mapping
and quantification

Transcriptome Mapping

Reads



*RSEM,
Kallisto,
Sailfish,
Salmon*

Transcript mapping
and quantification

Biological samples/Library preparation



Sequence reads



FASTQ

(+reference transcriptome index)

Quantify expression

Salmon, Kallisto, Sailfish



Pseudocounts

DGE with R



Functional Analysis
with R



Goal: Finding where in the genome these reads originated from

Genome

chrX: 52139280 152139290 152139300 152139310 152139320 152139330
--->CGCCGTCCCTCAGAAATGGAAACCTCGCTTCTCTCTGCCCCACAATGCGCAAGTCAG

Sequence reads

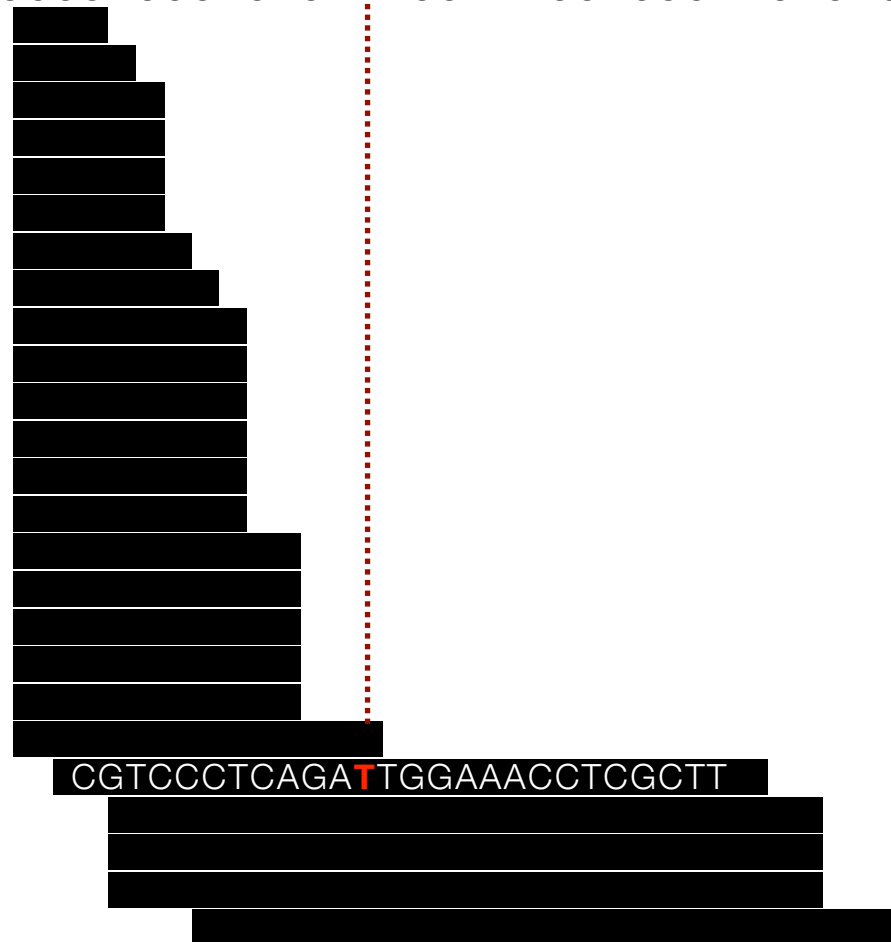
CGTCCCTCAGAAATGGAAACCTCGCTT

A simple case of string matching

Genome

chrX: 52139280 152139290 152139300 152139310 152139320 152139330
--->CGCCGTCCCTCAGAAATGGAAACCTCGCTTCTCTCTGCCCCACAATGCGCAAGTCAG

Sequence reads



A simple case of string matching?

Non-comprehensive list of challenges

- Large, incomplete and repetitive genomes OR transcriptomes with overlapping transcripts (isoforms)
- Short reads: 50-150 bp
 - Non-unique alignment
 - Sensitive to non-exact matching (variants, sequencing errors)
- Massive number of short reads
- Small insert size: 200-500 bp libraries
- Compute capacity for efficient mapping

Building an index

- Having an index of the reference sequence provides an efficient way to search
- Once index is built, it can be queried any number of times
- Every genome or transcriptome build requires a new index for the specific tool in question.


Commonly used indexing methods

- Hash-based (Salmon, Kallisto)
- Suffix arrays (Salmon, STAR)
- Burrows-Wheeler Transform (BWA, Bowtie2)

Hash-based alignment (circa 1990)



- ▶ Pick k-mer size, build lookup of every k-mer in the reference mapped to its positions (the index)
- ▶ Break the query into k-mers
- ▶ Seed-and-extend strategy
- ▶ For BLAST, 100% match the query k-mer to reference then extend until score drops below 50%
- ▶ 0.1 - 1 sec per query; not feasible for NGS data



Hash-based alignment (present day)

- ▶ Need to make some concessions on sensitivity by making adaptations for use on NGS data:
 - ▶ allow for mismatches and/or gaps (ELAND, MAQ, SOAP)
 - ▶ using multiple seeds (BLAT, ELAND2)
- ▶ Memory intensive and slower (~16GB RAM required for hg19)
- ▶ Simpler in design but more sensitive

Suffix arrays

- ▶ A sorted table of all suffixes (substrings) of a given string
- ▶ A suffix array will contain integers that represent the starting indexes of the all the suffixes of a given string, after the aforementioned suffixes are sorted
- ▶ Requires large amount of memory to load the suffix array and genome sequence prior to alignment
- ▶ Popular Tools:

STAR (2012), Salmon

Let the given string be “mississippi”

| Suffixes | ID | Sorted Suffixes | Suffix Array |
|---------------|----|-----------------|--------------|
| mississippi\$ | 1 | \$ | 12 |
| ississippi\$ | 2 | i\$ | 11 |
| ssissippi\$ | 3 | ippi\$ | 8 |
| sissippi\$ | 4 | issippi\$ | 5 |
| issippi\$ | 5 | ississippi\$ | 2 |
| ssippi\$ | 6 | mississippi\$ | 1 |
| sippi\$ | 7 | pi\$ | 10 |
| ippi\$ | 8 | ppi\$ | 9 |
| ppi\$ | 9 | sippi\$ | 7 |
| pi\$ | 10 | sissippi\$ | 4 |
| i\$ | 11 | ssippi\$ | 6 |
| \$ | 12 | ssissippi\$ | 3 |

The suffix array will be:
{12, 11, 8, 5, 2, 1, 10, 9, 7, 4, 6, 3}

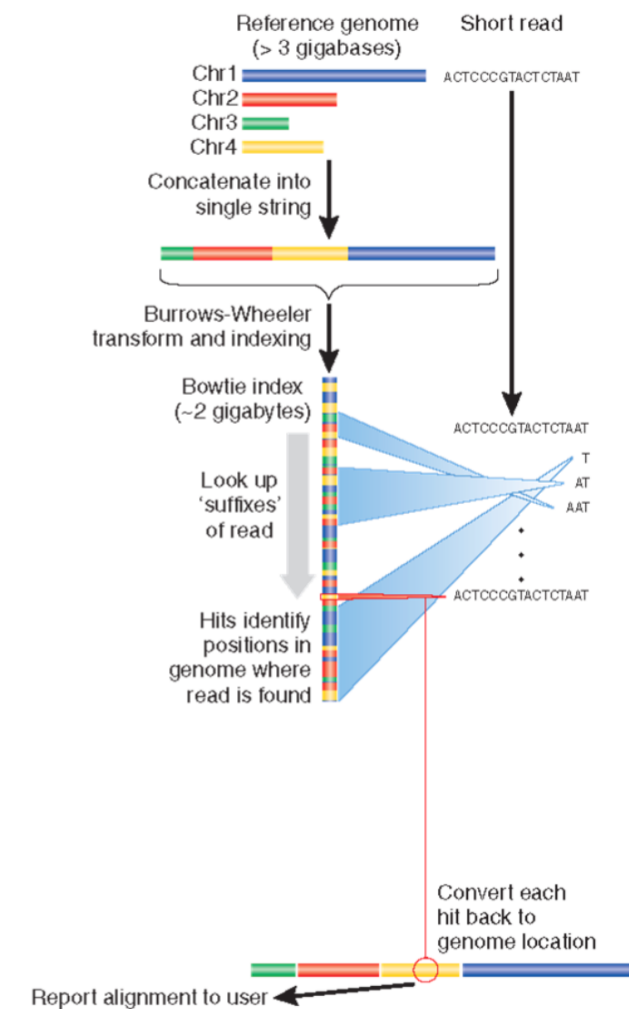
Burrows-Wheeler transform

- ▶ A compressed form of suffix arrays
- ▶ Tends to put runs of the same character together rather than alphabetically, which makes the compression work well

| Suffixes | ID | Sorted Suffixes | Suffix Array | Sorted Rotations (A_s matrix) | BWT Output (L) |
|---------------|----|-----------------|--------------|----------------------------------|--------------------|
| mississippi\$ | 1 | \$ | 12 | \$mississippi | i |
| ississippi\$ | 2 | i\$ | 11 | i\$mississipp | p |
| ssissippi\$ | 3 | ippi\$ | 8 | ippi\$mississ | s |
| sissippi\$ | 4 | issippi\$ | 5 | issippi\$miss | s |
| issippi\$ | 5 | ississippi\$ | 2 | ississippi\$m | m |
| ssippi\$ | 6 | mississippi\$ | 1 | mississippi\$ | \$ |
| sippi\$ | 7 | pi\$ | 10 | pi\$mississip | p |
| ippi\$ | 8 | ppi\$ | 9 | ppi\$mississi | i |
| ppi\$ | 9 | sippi\$ | 7 | sippi\$missis | s |
| pi\$ | 10 | sissippi\$ | 4 | sissippi\$mis | s |
| i\$ | 11 | ssippi\$ | 6 | ssippi\$missi | i |
| \$ | 12 | ssissippi\$ | 3 | ssissippi\$mi | i |

Burrows-Wheeler transform

- ▶ Much less memory because of compression;
~1.5 GB of RAM required for hg19 index
- ▶ But compression results in diminished efficiency of the string search operations
- ▶ Popular Tools:
 - Bowtie2 (2012)
 - SOAP2
 - BWA-MEM (2013)



Reference data versions matter

- Ensembl, UCSC and NCBI all often use the same genome assemblies or builds (e.g. GrCh38 == hg38)
- Make sure that the annotation file (GTF) is exactly matched with the genome file (fasta), or transcriptome file (fasta)
 - Same build version
 - Same source (e.g. both from FlyBase)

How does Salmon map reads?

Lightweight alignment and quantification using Salmon




Image source: [RNA-seq blog](#)

Lightweight alignment and quantification using Salmon




Image source: [RNA-seq blog](#)

Lightweight alignment and quantification using Salmon

- ▶ **Reference:** FASTA file of all transcript sequences for the organism

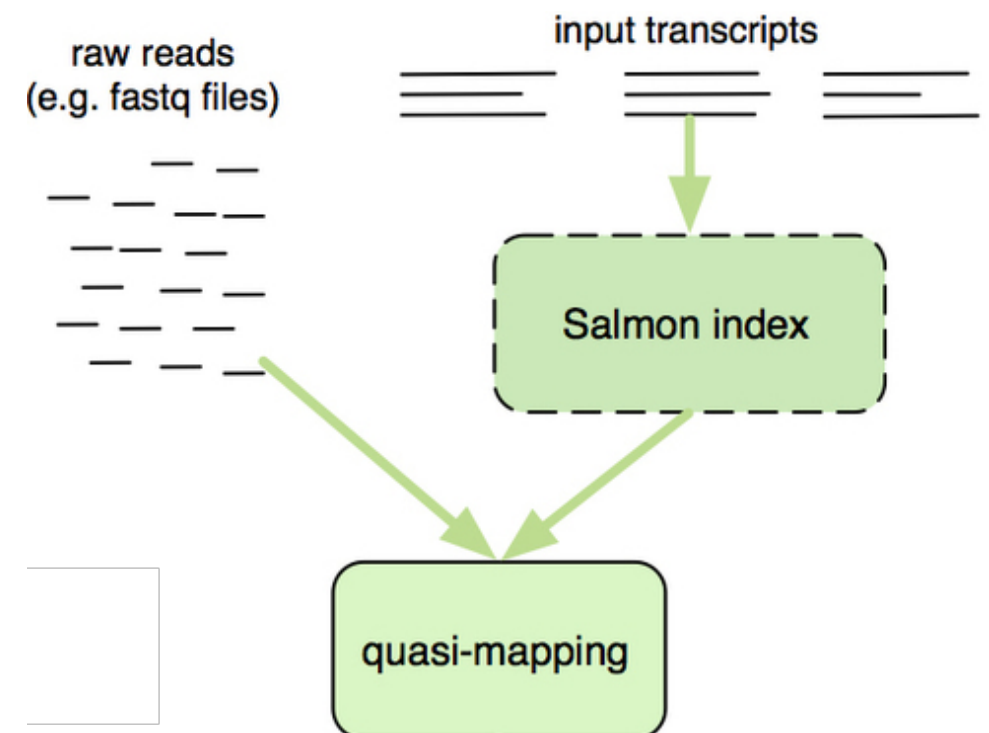


Image source: [RNA-seq blog](#)

Lightweight alignment and quantification using Salmon

- ▶ **Reference:** FASTA file of all transcript sequences for the organism
- ▶ **Index:** (2 components)
 - ▶ Suffix array
 - ▶ Hash table (mapping each transcript to its location in the SA)




Image source: [RNA-seq blog](#)

Lightweight alignment and quantification using Salmon

- ▶ **Reference:** FASTA file of all transcript sequences for the organism
- ▶ **Index:** (2 components)
 - ▶ Suffix array
 - ▶ Hash table (mapping each transcript to its location in the SA)
- ▶ **Output:** abundance estimates
 - ▶ Number of reads mapping to each transcript listed in the reference





Image source: [RNA-seq blog](#)

Lightweight alignment and quantification using Salmon

- ▶ **Reference:** FASTA file of all transcript sequences for the organism
- ▶ **Index:** (2 components)
 - ▶ Suffix array
 - ▶ Hash table (mapping each transcript to its location in the SA)
- ▶ **Output:** abundance estimates
 - ▶ Number of reads mapping to each transcript listed in the reference




Note that we don't get any information on *where* each read is mapping!

Why use lightweight alignment?

- Approaches avoid base-to-base alignment
- Faster, more efficient (~ >20x faster than alignment-based)
- Improved accuracy for transcript-level quantification
- Improvements in accuracy for gene-level quantification**
- Other tools include: Kallisto (quasi-aligner), Sailfish (kmer-based), Salmon (quasi-aligner), RSEM

**doi: [10.12688/f1000research.7563.2](https://doi.org/10.12688/f1000research.7563.2)



Quality Checks

QC metrics

Various metrics can give important information about the quality of the library:

- Total % of reads aligning? % of uniquely mapping reads? % of properly paired PE reads?
- Genomic origin of reads (exonic, intronic, intergenic)
- Quantity of rRNA
- Transcript coverage and 5'-3' bias

How do we compute QC metrics?

- Tools like RNA-SeQC and Qualimap
 - *Input:* alignment file formats (i.e. SAM/BAM)
 - *Output:* summary of the different metrics in an HTML report format


SAM/BAM file format

- Sequence Alignment Map (SAM) format contains information on a per-read basis:
 - Coordinates of alignment, including strand
 - Mismatches
 - Mapping information (unique?, properly paired?, etc.)
 - Quality of mapping (tool-specific scoring systems)
- BAM: Binary version of SAM alignment format files



[More information about SAM/BAM](#)

Where do we get this SAM/BAM file from?




- Need to *align reads to the genome*
- Genome alignment tools perform base-to-base alignment of reads to the genome sequence (FASTA) and can be guided with a gene annotation file (GTF)
- Genome alignment outputs a SAM/BAM file



Splice-aware alignment



Splice-aware alignment



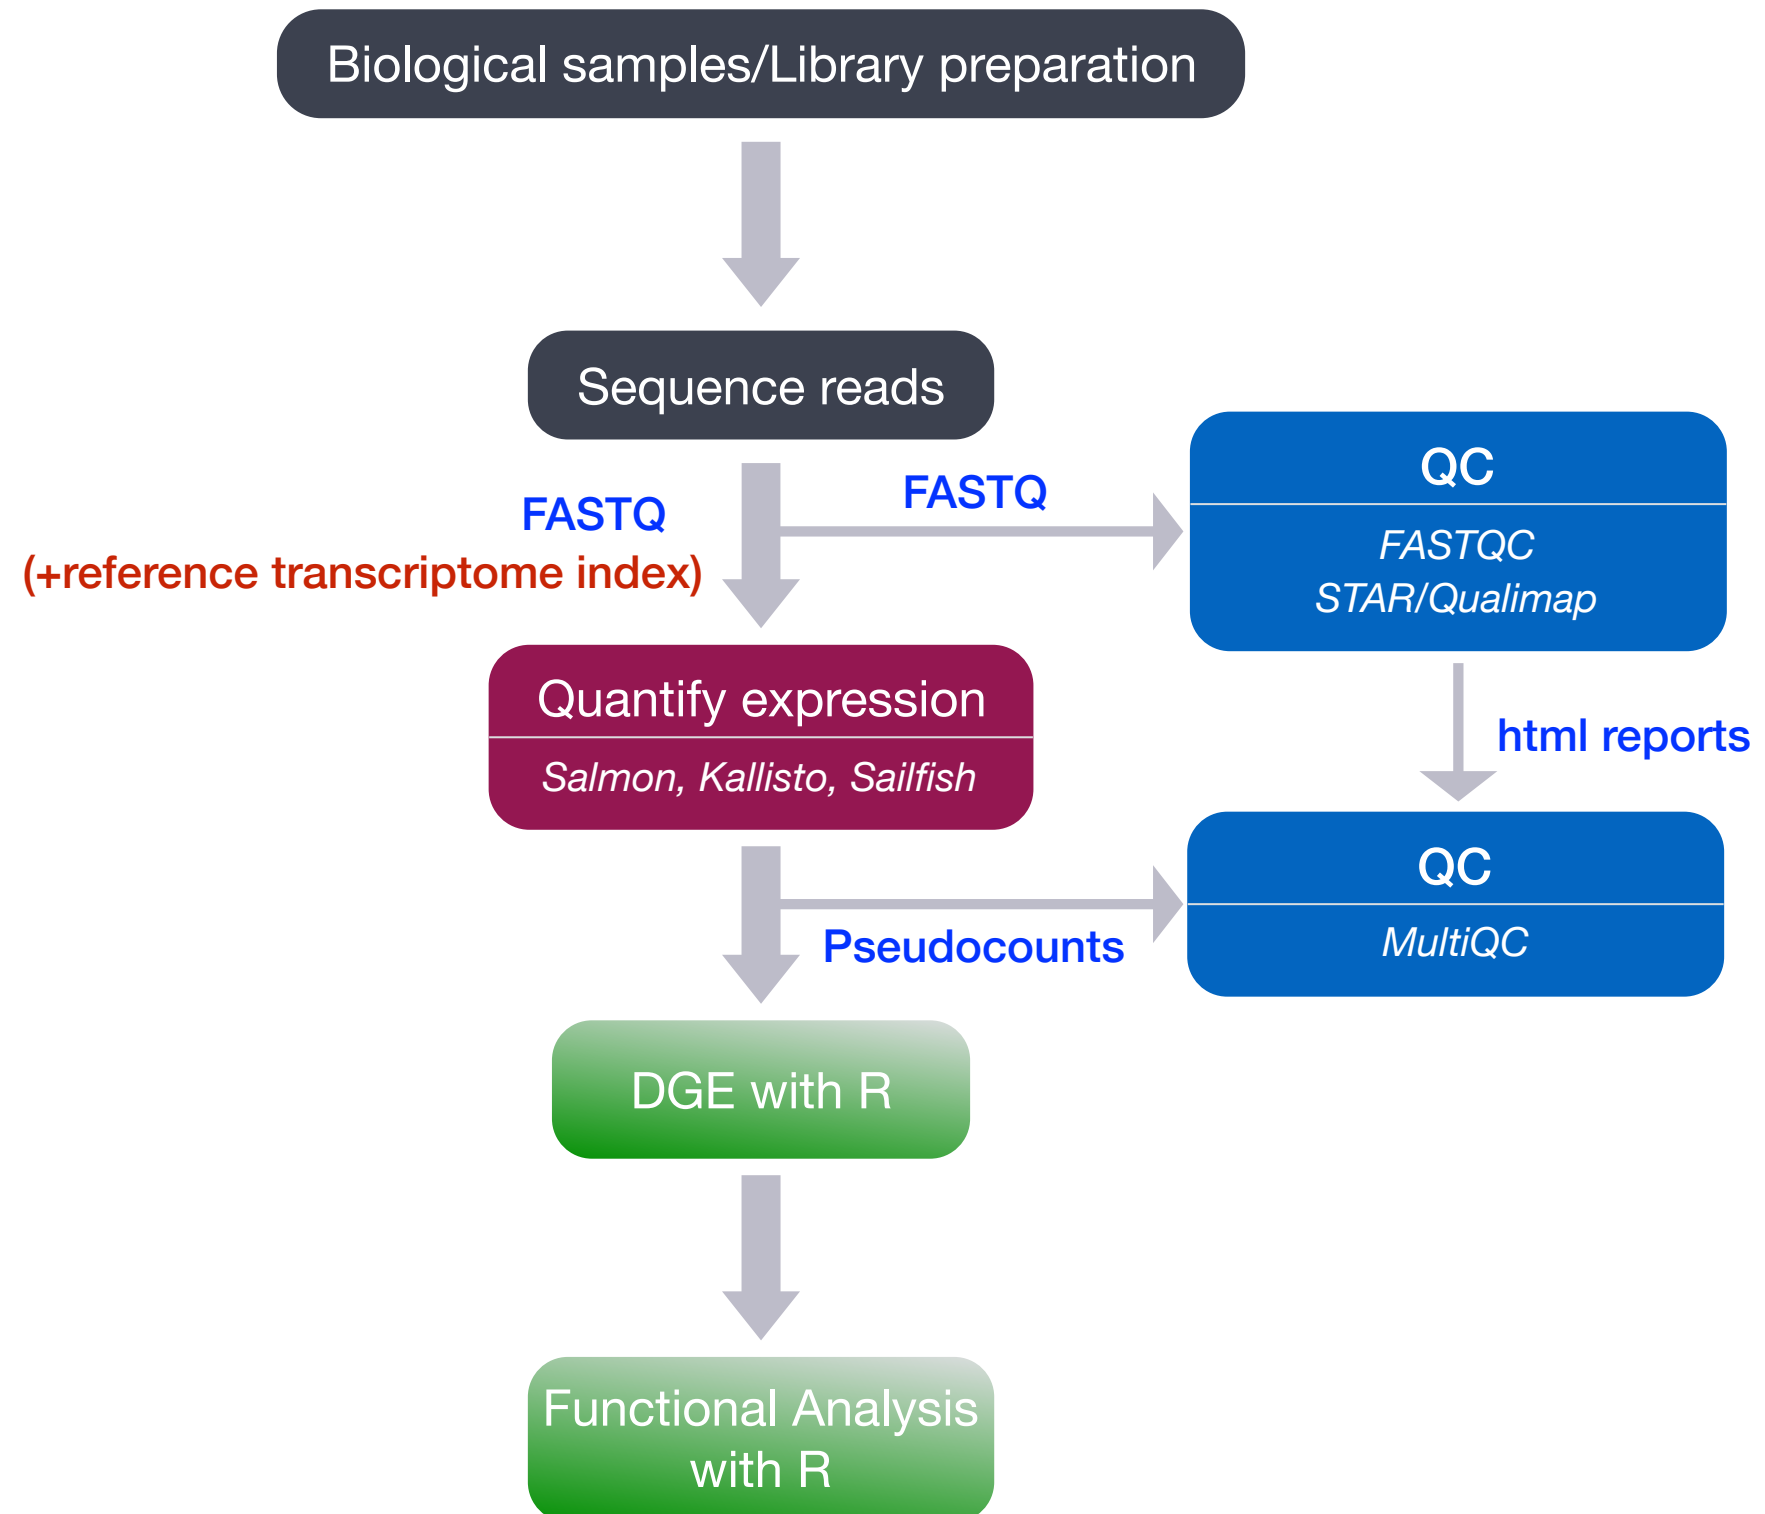
Splice-aware alignment

Splice-aware alignment tools:

HISAT2, STAR, MapSplice, SOAPSplice, Passion, SpliceMap,
RUM, ABMapper, CRAC, GSNAP, HMMSplicer, Olego, BLAT

There are excellent genome aligners available that are not splice-aware. These are useful for aligning directly to genes. However, you will lose isoform information.

Bowtie2, BWA, Novoalign (not free), SOAPaligner



These materials have been developed by members of the teaching team at the Harvard Chan Bioinformatics Core (HBC). These are open access materials distributed under the terms of the Creative Commons Attribution license (CC BY 4.0), which permits unrestricted use, distribution, and reproduction in any medium, provided the original author and source are credited.

